



香港中文大學(深圳)
The Chinese University of Hong Kong, Shenzhen

Introduction to Computer Science: Programming Methodology

Lecture 8 Data Structure and Algorithm - Intro

**Prof. Guiliang Liu
School of Data Science**

Data structure and algorithm

- A **data structure** is a systematic way of organizing and accessing data
- An **algorithm** is a step-by-step procedure for performing some task in a finite amount of time.

Why study data structure and algorithm?

- Important for **all other branches** of computer science
- Plays a **key role** in modern technological innovation
- **Moore's law** predicts that the density of transistors in integrated circuits would continue to double every 1 to 2 years
- However, in many areas, performance gains due to the **improvements in algorithms** have **greatly exceeded** even the dramatic performance gains due to increased processor speed

Why study data structure and algorithm?

- Provide novel “lens” on processes outside of computer science and technology, such as quantum mechanics, economic markets, evolution
- Challenging (good for your brain!!) and funny

Example: Integer Multiplication

- **Inputs:** two n-digits number x and y
- **Output:** the product of x and y
- **Primitive operations:** add or multiply 2 single digit numbers

The algorithm designer' s mantra

- “Perhaps the most important principle for the good algorithm designer is to refuse to be content”

Aho, Hopcroft, and Ullman, *The Design and Analysis of Computer Algorithms*, 1974

How do we define a “good” algorithm?

- The primary analysis of algorithms involves characterizing the **running times** and **space usage** of algorithms and data structure operations
- **Running time** is a natural measure of “goodness,” since time is a precious resource—computer solutions should run as fast as possible
- **Space usage** is another major issue to consider when we design an algorithm, since we only have limited storage spaces

Measuring the running time experimentally

```
from time import time

startTime = time()

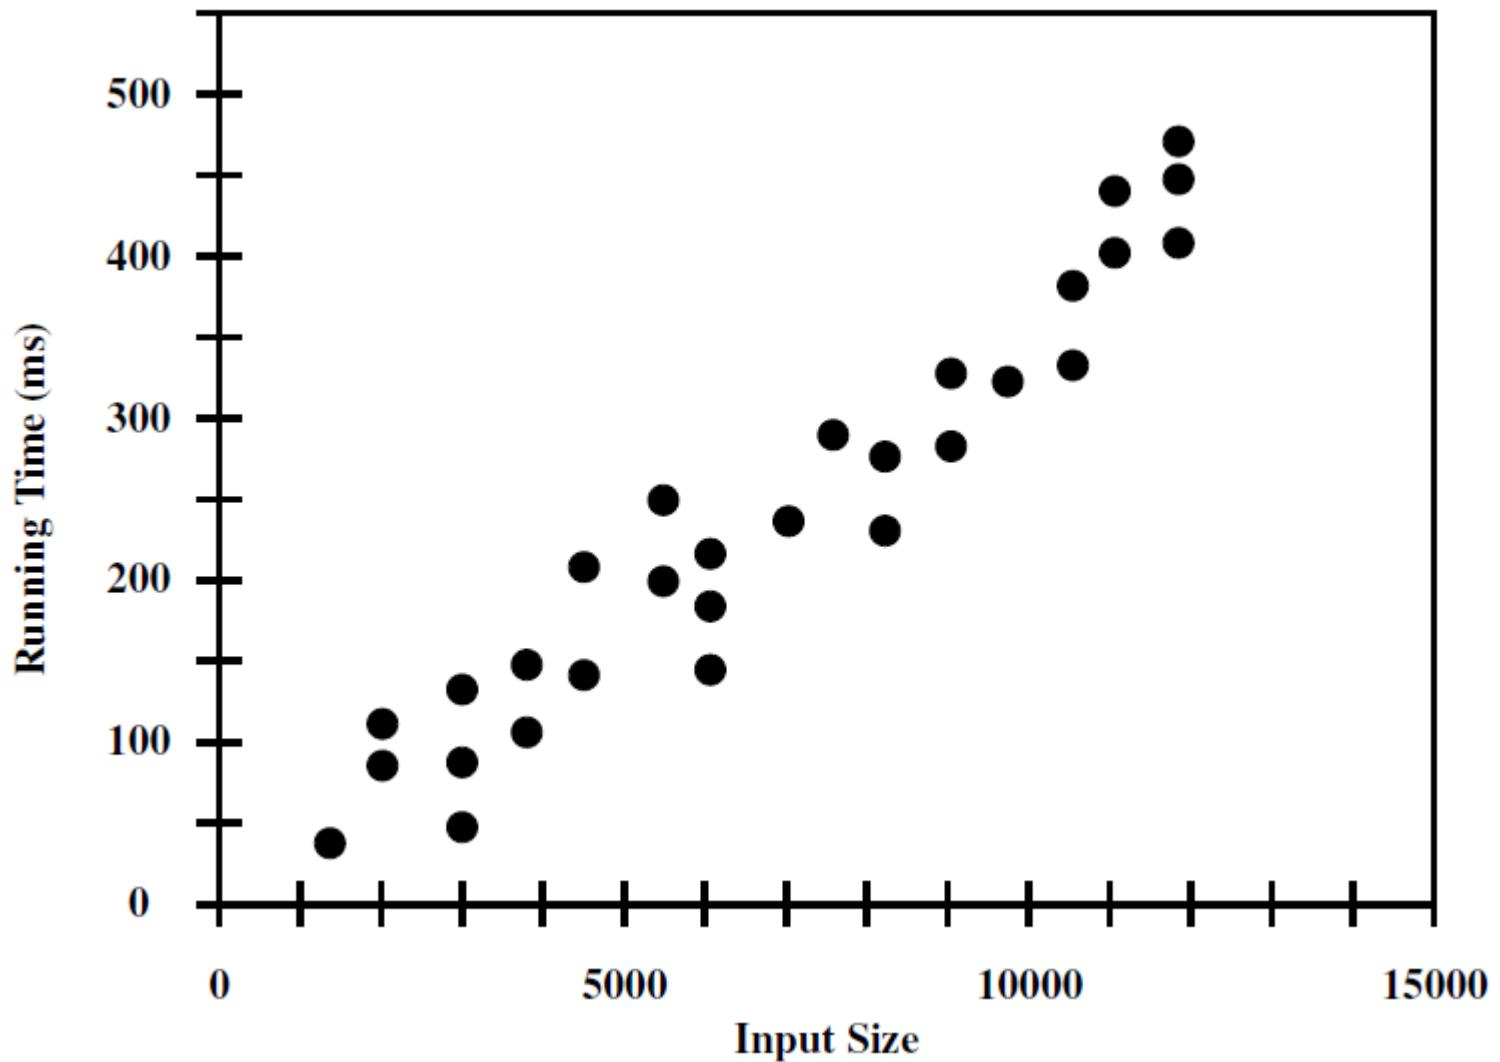
for i in range(1, 20000):
    if i%10 == 0:
        print(i)

endTime = time()

print('The time elapsed is:', endTime - startTime, 'seconds')
```

Visualize the running time

- Running time and space usage are dependent on the **size of the input**
- Perform independent experiments on many different **test inputs of various sizes**
- Visualize the results by plotting the performance of each run of the algorithm as a point



Challenges of experimental analysis

- Experimental running times of two algorithms are difficult to directly compare unless the experiments are performed in the same hardware and software environments
- Experiments can be done only on a limited set of test inputs; hence, they leave out the running times of inputs not included in the experiment (and these inputs may be important)
- An algorithm must be fully implemented in order to execute it to study its running time experimentally

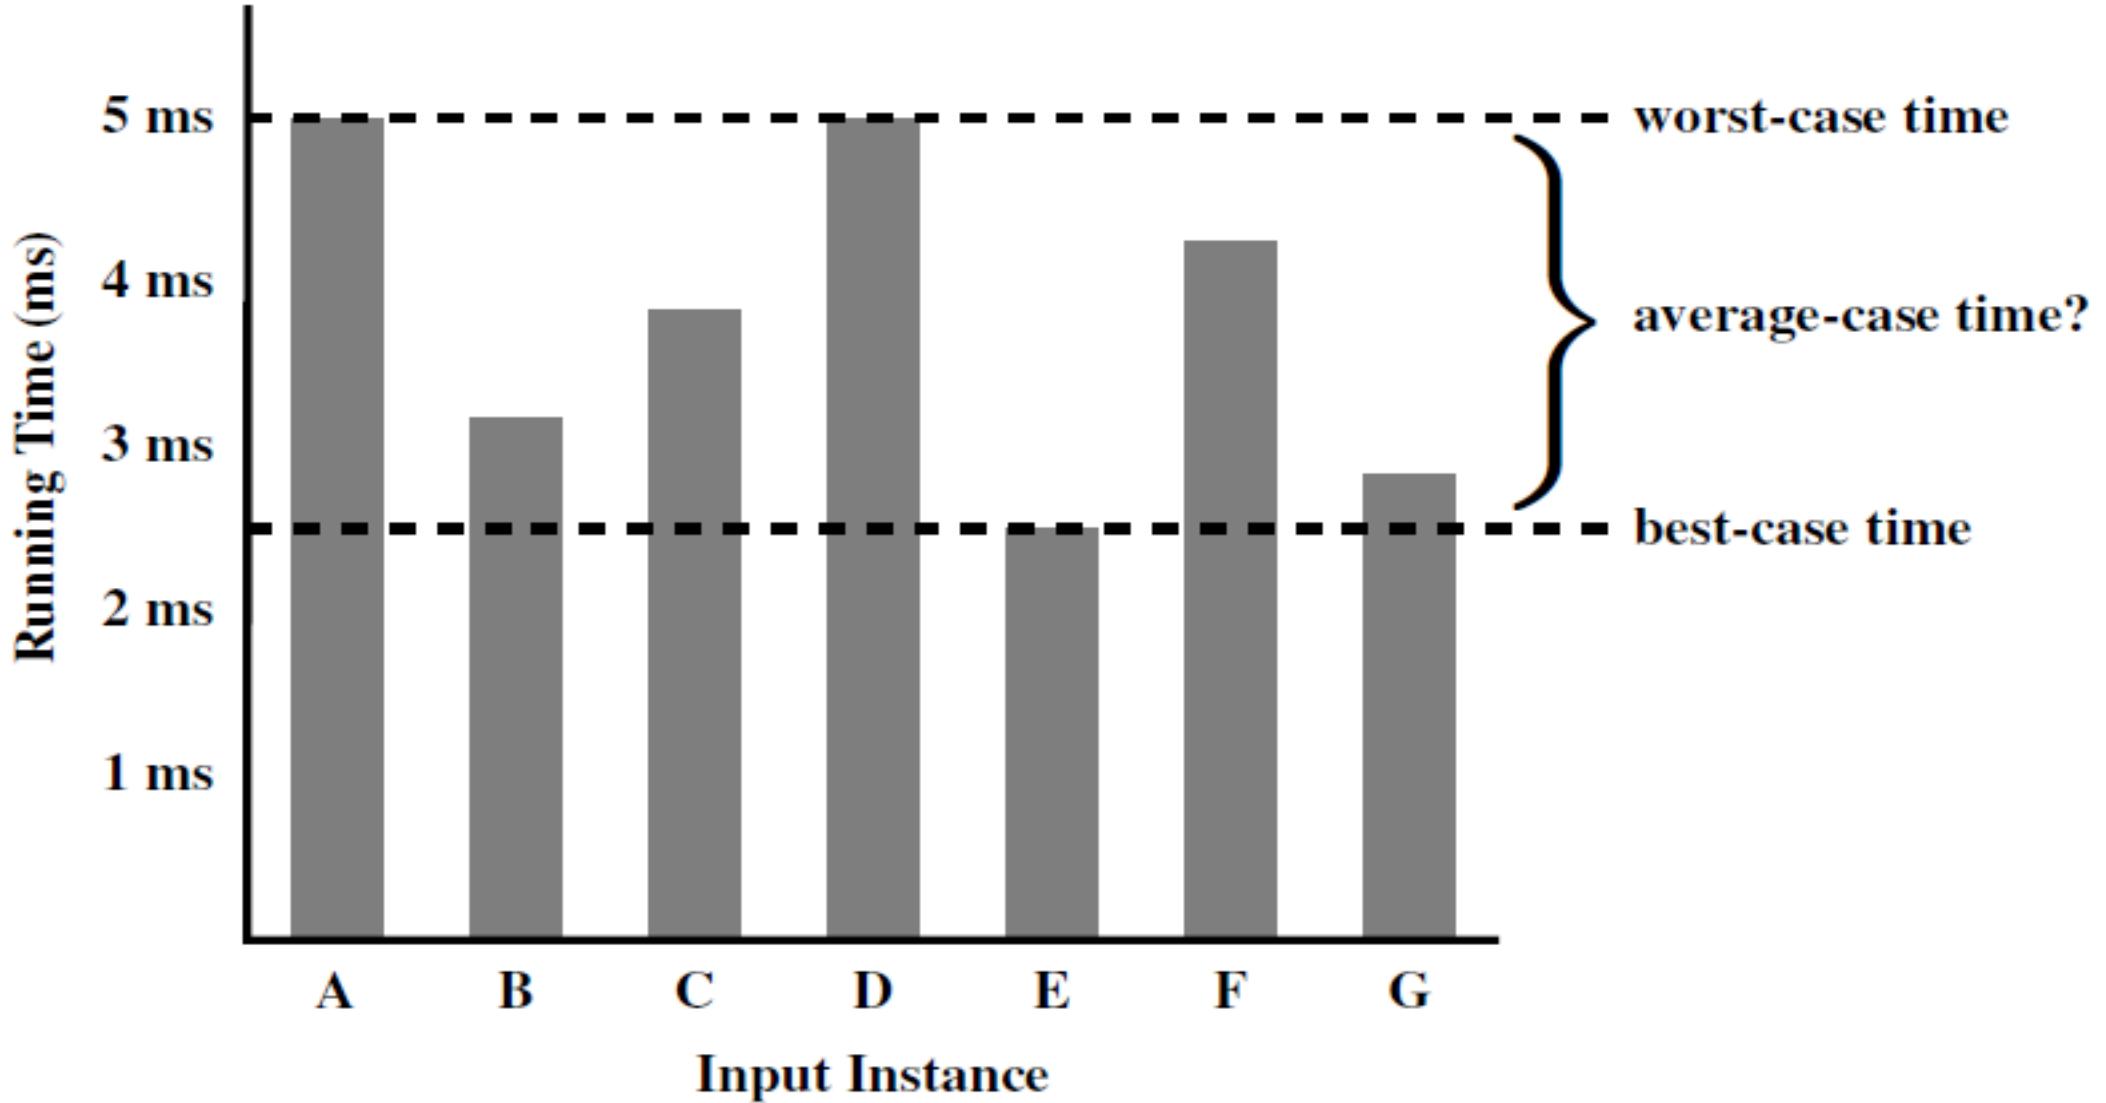
Principle of algorithm analysis 1: Counting primitive operations

- To analyse the running time of an algorithm without performing experiments, we perform an analysis directly on a high-level description of the algorithm
- We define a set of primitive operations such as the following:
 - ✓ Assigning an identifier to an object
 - ✓ Determining the object associated with an identifier
 - ✓ Performing an arithmetic operation (for example, adding two numbers)
 - ✓ Comparing two numbers
 - ✓ Accessing a single element of a Python list by index
 - ✓ Calling a function (excluding operations executed within the function)
 - ✓ Returning from a function.

Principle of algorithm analysis 2: Measuring Operations as a Function of Input Size


- To capture the order of growth of an algorithm's running time, we will associate, with each algorithm, a function $f(n)$ that characterizes the number of primitive operations that are performed as a function of the input size n

Principle of algorithm analysis 3: Focusing on the Worst-Case Input



The 7 functions used in algorithm analysis

- We may use the following 7 functions to measure the time complexity of an algorithm: **constant, logarithm, linear, N-log-N, quadratic, cubic and other polynomials, exponential**



Asymptotic Analysis

- In algorithm analysis, we focus on the growth rate of the running time as a function of the input size n , taking a “big-picture” approach
- Vocabulary for the analysis and design of algorithms
- “Sweet spot” for high-level reasoning about algorithms
- Coarse enough to suppress unnecessary details, e.g. architecture/language/compiler...
- Sharp enough to make meaningful comparisons between algorithms

The big Oh notation

- Let $f(n)$ and $g(n)$ be functions mapping positive integers to positive real numbers.
- We say that $f(n)$ is $O(g(n))$ if there is a real constant $c > 0$ and an integer constant $n_0 \geq 1$ such that

$$f(n) \leq cg(n), \text{ for } n \geq n_0$$

- This definition is often referred to as the “big-Oh” notation
- **Example:** The function $8n+5$ is $O(n)$.

The big Oh notation

- The big-Oh notation allows us to say that a function $f(n)$ is “less than or equal to” another function $g(n)$ up to a constant factor and **in the asymptotic sense** as n grows toward infinity
- The big-Oh notation is used widely to characterize **running times** and **space bounds** in terms of some parameter **n** , which varies from problem to problem, but is always defined as a chosen measure of the “**size**” of the problem

Some Properties of the Big-Oh Notation

- The big-Oh notation allows us to ignore constant factors and lower-order terms and focus on the main components of a function that affect its growth
- Example: $5n^4 + 3n^3 + 2n^2 + 4n + 1$ is $O(n^4)$
- Example: 2^{n+2} is $O(2^n)$
- Example: $2n + 100\log n$ is $O(n)$
- In general, we should use the big-Oh notation to characterize a function as closely as possible

Comparative analysis

Question: Suppose two algorithms solving the same problem are available: an algorithm A, which has a running time of $O(n)$, and an algorithm B, which has a running time of $O(n^2)$. Which algorithm is better?

Answer: Algorithm A is **asymptotically better** than algorithm B

Comparative analysis

- We can use the big-Oh notation to order classes of functions by asymptotic growth rate
- Our seven functions are ordered by increasing growth rate in the following sequence

n	$\log n$	n	$n \log n$	n^2	n^3	2^n
8	3	8	24	64	512	256
16	4	16	64	256	4,096	65,536
32	5	32	160	1,024	32,768	4,294,967,296
64	6	64	384	4,096	262,144	1.84×10^{19}
128	7	128	896	16,384	2,097,152	3.40×10^{38}
256	8	256	2,048	65,536	16,777,216	1.15×10^{77}
512	9	512	4,608	262,144	134,217,728	1.34×10^{154}

The line of tractability

- To differentiate **efficient** and **inefficient** algorithms, the general line is between **polynomial time algorithms** and **exponential time algorithms**
- The distinction between polynomial-time and exponential-time algorithms is considered a robust measure of **tractability**

Example: Finding the smallest number in a list

```
smallest_so_far = None
print('Before', smallest_so_far)

for num in [9, 39, 21, 98, 4, 5, 100, 65]:
    if smallest_so_far == None:
        smallest_so_far = num
    elif num < smallest_so_far:
        smallest_so_far = num
    print(smallest_so_far, num)

print('After', smallest_so_far)
```

- What is the time complexity of this algorithm?

Recursion

- Recursion is a technique by which a function makes one or more calls to itself during execution
- Recursion provides an elegant and powerful alternative for performing repetitive tasks
- Recursion is an important technique in the study of data structures and algorithms

Inception



Example: The factorial function

- The **factorial** of a positive integer n , denoted $n!$, is defined as follows:

$$n! = \begin{cases} 1 & \text{if } n = 0 \\ n \cdot (n - 1) \cdot (n - 2) \cdots 3 \cdot 2 \cdot 1 & \text{if } n \geq 1. \end{cases}$$

- The factorial function is important because it is known to equal the number of ways in which n distinct items can be arranged into a sequence, that is, the number of permutations of n items

The recursive definition

- First, a recursive definition contains one or more **base cases**, which are defined **non-recursively** in terms of fixed quantities
- Second, it also contains one or more **recursive cases**, which are defined by appealing to the definition of the function being defined

The recursive definition of factorial function

- The factorial function can be naturally defined in a recursive way, for example, $5! = 5 \cdot (4 \cdot 3 \cdot 2 \cdot 1) = 5 \cdot 4!$
- More generally, for a positive integer n , we can define $n!$ to be $n \cdot (n-1)!$
- Therefore, the recursive definition of factorial function is:

$$n! = \begin{cases} 1 & \text{if } n = 0 \\ n \cdot (n-1)! & \text{if } n \geq 1. \end{cases}$$


Solution

```
def facFunc(n):  
  
    if n<0:  
        print(' Invalid input. ')  
        return None  
    elif n == 0:  
        return 1  
    else:  
        return n*facFunc(n-1)
```

How Python implements recursion

- In Python, each time a function (recursive or otherwise) is called, a structure known as an **activation record** or **frame** is created to store information about the progress of that invocation of the function
- This activation record stores the function call's **parameters** and **local variables**
- When the execution of a function leads to a nested function call, the execution of the former call is suspended and its activation record stores the place in the source code at which the **flow of control should continue** upon return of the nested call

The recursive trace



Worst-case time complexity? $O(n)$

Example: Drawing an English ruler

- We denote the length of the tick designating a whole inch as the ***major tick length***.
- Between the marks for whole inches, the ruler contains a series of ***minor ticks***, placed at intervals of $1/2$ inch, $1/4$ inch, and so on.
- As the size of the interval decreases by half, the tick length decreases by one

----- 0

-

--

-

-

--

-

----- 1

-

--

-

-

--

-

----- 2

-

--

-

----- 0

-

--

-

-

--

-

----- 1

-

--

-

-

--

-

----- 2

-

--

-

(a)

----- 1

(b)

----- 0

-

--

-

----- 1

-

--

-

----- 2

-

--

-

----- 3

-

--

-

(c)

Recursive implementation of English ruler

- An interval with a central tick length $L \geq 1$ is composed of:
 - ✓ An interval with a central tick length $L-1$
 - ✓ A single tick of length L
 - ✓ An interval with a central tick length $L-1$

Solution

```
def draw_line(tickLen, tickLabel=' '):
    line = '-' * tickLen
    if tickLabel:
        line += ' ' + tickLabel
    print(line)


def draw_interval(centerLen):
    if centerLen > 0:
        draw_interval(centerLen - 1)
        draw_line(centerLen)
        draw_interval(centerLen - 1)

def draw_ruler(numInch, majorLen):
    draw_line(majorLen, '0')

    for j in range(1, 1 + numInch):
        draw_interval(majorLen - 1)
        draw_line(majorLen, str(j))
```

The recursive trace for English ruler

- Worst-case time complexity? $O(2^L)$
- L denotes the major tick length



Example: Binary search

- A classic and very useful recursive algorithm, **binary search**, can be used to efficiently locate a target value within a **sorted** sequence of **n** elements
- When the sequence is **unsorted**, the standard approach to search for a target value is to use a loop to examine every element, until either finding the target or exhausting the data set; This is known as the **sequential search algorithm**

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
2	4	5	7	8	9	12	14	17	19	22	25	27	28	33	37

Binary search

- When the sequence is sorted and indexable, binary search is a much more efficient algorithm
- For any index j , we know that all the values stored at indices $0, \dots, j-1$ are less than or equal to the value at index j , and all the values stored at indices $j+1, \dots, n-1$ are greater than or equal to that at index j

The strategy of binary search

- We call an element of the sequence a **candidate** if, at the current stage of the search, we cannot rule out that this item matches the target
- The algorithm maintains two parameters, **low** and **high**, such that all the candidate entries have index at least **low** and at most **high**
- Initially, **low = 0** and **high = n-1**. We then compare the target value to the median candidate, that is, the item **data[mid]** with index
$$\text{mid} = \lfloor (\text{low}+\text{high})/2 \rfloor$$

The strategy of binary search

- If the target equals `data[mid]`, then we have found the item we are looking for, and the search terminates successfully
- If `target < data[mid]`, then we recur on the first half of the sequence, that is, on the interval of indices from `low` to `mid-1`
- If `target > data[mid]`, then we recur on the second half of the sequence, that is, on the interval of indices from `mid+1` to `high`

Solution

```
def binarySearch(data, target, low, high):
    if low>high:
        print(' Cannot find the target number! ')
        return False
    else:
        mid = (low+high)//2
        if target==data[mid]:
            print(' The target number is at position ', mid)
            return True
        elif target<data[mid]:
            return binarySearch(data, target, low, mid-1)
        else:
            return binarySearch(data, target, mid+1, high)

def main():
    data = [1, 3, 5, 6, 16, 78, 100, 135, 900]
    target = 16
    binarySearch(data, target, 0, len(data)-1)
```

Time complexity of binary search

Proposition: The binary search algorithm runs in $O(\log n)$ time for a sorted sequence with n elements

Why ?

Proof

Justification: To prove this claim, a crucial fact is that with each recursive call the number of candidate entries still to be searched is given by the value $\text{high} - \text{low} + 1$.

Moreover, the number of remaining candidates is reduced by at least one half with each recursive call. Specifically, from the definition of mid , the number of remaining candidates is either

$$(\text{mid} - 1) - \text{low} + 1 = \left\lfloor \frac{\text{low} + \text{high}}{2} \right\rfloor - \text{low} \leq \frac{\text{high} - \text{low} + 1}{2}$$

or

$$\text{high} - (\text{mid} + 1) + 1 = \text{high} - \left\lfloor \frac{\text{low} + \text{high}}{2} \right\rfloor \leq \frac{\text{high} - \text{low} + 1}{2}.$$

Initially, the number of candidates is n ; after the first call in a binary search, it is at most $n/2$; after the second call, it is at most $n/4$; and so on. In general, after the j^{th} call in a binary search, the number of candidate entries remaining is at most $n/2^j$. In the worst case (an unsuccessful search), the recursive calls stop when there are no more candidate entries. Hence, the maximum number of recursive calls performed, is the smallest integer r such that

$$\frac{n}{2^r} < 1.$$

In other words (recalling that we omit a logarithm's base when it is 2), $r > \log n$. Thus, we have $r = \lfloor \log n \rfloor + 1$,

which implies that binary search runs in $O(\log n)$ time. ■