

# Lecture 17 - Policy gradient

Guiliang Liu

The Chinese University of Hong Kong, Shenzhen

DDA4230: Reinforcement Learning

Course Page: [\[Click\]](#)

# Policy-based and value-based algorithms

**Value-based algorithms** include Q-learning, temporal-difference learning, and policy and value iteration

- These algorithms learn the values of actions  $V(s)$  or  $Q(s, a)$  and then selected action  $a$  based on the action values  $\pi(s) = \arg \max_{a \in \mathcal{A}} Q(s, a)$ ;
- The policy does not exist without the action value estimates  $Q(s)$ .



香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen

# Policy-based and value-based algorithms

Concerns about value-based methods.

- The vanilla approaches can only address discrete action spaces due to the  $\arg \max_{a \in \mathcal{A}}$  operation. However, in practice, the action space is usually continuous.
- Computing the action value functions  $Q(s, a)$  for all state-action pair is costly when the action and state spaces are large or continuous.
- The policy of Q-Learning is deterministic and  $\varepsilon$ -greedy explore can be inefficient.
- It implicitly and indirectly improves the policy by improving the estimates of the values functions. However, we would think intuitively that improving the policy directly would be more efficient.



香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen

# Policy-based and value-based algorithms

**Policy gradient** is the canonical approach for policy-based learning.

- Policy-based method directly parameterizes the policy function  $\pi_\theta(s)$  without calculating the value functions.
- We use the notation  $\theta \in R^d$  for the policy's parameter vector. We then write  $\pi(a | s, \theta) = \mathbb{P}(a_t = a | s_t = s, \theta)$  as the probability that action  $a$  is taken given that the environment is in state  $s$  with parameter  $\theta$ .
- A value function may still be used to **learn** the policy parameter, but is not required for action selection (will talk about it later in the actor-critic algorithm).



香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen

# Policy approximation with parametrization

**Discrete Action Space.** then a natural way to parameterize a policy is to form parameterized state-action preferences  $h(s, a, \theta)$  for each  $(s, a)$  pair and use a softmax distribution

$$\pi(a | s, \theta) = \frac{\exp(h(a, s, \theta))}{\sum_{a'} \exp(h(a', s, \theta))}. \quad (\text{softmax in action preferences})$$

- The state-action preference measures how the policy  $\pi_\theta$  prefer action  $a$  given state  $s$ . The actions with the highest preferences in each state are given the highest probabilities of being selected.



香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen

# Policy approximation with parametrization

**Discrete Action Space.** then a natural way to parameterize a policy is to form parameterized state-action preferences  $h(s, a, \theta)$  for each  $(s, a)$  pair and use a softmax distribution

$$\pi(a | s, \theta) = \frac{\exp(h(a, s, \theta))}{\sum_{a'} \exp(h(a', s, \theta))}. \quad (\text{softmax in action preferences})$$

- The action preferences  $h(a, s, \theta)$  can be parameterized arbitrarily. For example, it can simply be the linear combinations of features (as for the feature vectors  $x(a, s)$ )

$$h(a, s, \theta) = \theta^T x(a, s).$$



香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen

# Policy approximation with parametrization

**Continuous Action Space.** The policy can be defined as the normal probability density over a real-valued scalar action, with mean and standard deviation given by parametric function approximators

$$\pi(a | s, \theta) = \frac{1}{\sigma(s, \theta_\sigma)\sqrt{2\pi}} \exp\left(-\frac{(a - \mu(s, \theta_\mu))^2}{2\sigma(s, \theta_\sigma)^2}\right).$$

- We divide the policy's parameter vector into two parts,  $\theta = [\theta_\mu, \theta_\sigma]$ .
- One possible way to parametrize the mean and standard deviation is

$$\mu(s, \theta) = \theta_\mu^T x_\mu(s), \quad \sigma(s, \theta) = \exp(\theta_\sigma^T x_\sigma(s)),$$

where  $x_\sigma(s)$  and  $x_\mu(s)$  are feature vectors.



香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen

# Policy approximation with parametrization

## Advantages of using parametrization

- It handles both discrete and continuous action spaces.
- It could be deterministic or stochastic. If the optimal policy is deterministic, then the preference values  $h(a, s, \theta)$  will be driven infinitely higher than all other actions.
- The choice of policy parametrization is sometimes a good way of injecting prior knowledge about the desired form of the policy into the learning.
- Policy gradient has stronger convergence guarantees than value-based method because of the smooth change in the probability.



香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen

# Policy Gradient Objective

Recall the gradient descent algorithm,  $\theta_{t+1} = \theta_t + \alpha \widehat{\nabla J(\theta)}$  where  $J(\theta)$  is our objective function and  $\alpha$  is the learning rate.

For the episodic case, where the episode terminates at some terminal state set, we define the objective function  $J(\theta)$  as

$$J(\theta) = V^{\pi_\theta}(s_0) = \sum_{s \in \mathcal{S}} \rho^{\pi_\theta}(s | s_0) r(s),$$

where  $s_0$  is the starting state,  $V^{\pi_\theta}(s_0)$  is the value function for  $\pi_\theta$ , and  $r(s) = \mathbb{E}_{a \sim \pi}[\mathcal{R}(s, a)]$  is the expected reward at  $s$  following  $\pi$ . The occupancy measure  $\rho^{\pi_\theta}(s | s_0) = \frac{1}{T} \sum_{t=0}^T \mathbb{P}(s_t = s | s_0, \pi_\theta)$ , where  $T$  is a random variable denoting the index of the terminal step.



香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen

# Policy Gradient Objective

For the continuing case, where the process continues infinitely, we define the objective function  $J(\theta)$  as the averaged reward over the time steps.

$$\begin{aligned} J(\theta) &= \lim_{T \rightarrow \infty} \frac{1}{T} \sum_{t=1}^T \mathbb{E}[r_t | s_0, \pi_\theta] \\ &= \lim_{t \rightarrow \infty} \mathbb{E}[r_t | s_0, \pi_\theta] \\ &= \sum_s \rho^{\pi_\theta}(s | s_0) r(s) \\ &= V^{\pi_\theta}(s_0), \end{aligned}$$

where the occupancy measure  $\rho^{\pi_\theta}(s | s_0) = \lim_{t \rightarrow \infty} \mathbb{P}(s_t = s | s_0, \pi_\theta)$  is the stationary distribution of the Markov chain under policy  $\pi_\theta$ .



香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen

# Policy Gradient Objective

For the discounted case where  $\gamma < 1$ , we define the objective function  $J(\theta)$  as the expected discounted return

$$J(\theta) = V^{\pi_\theta}(s_0) = \sum_{s \in \mathcal{S}} \rho^{\pi_\theta}(s | s_0) r(s),$$

where the occupancy measure  $\rho^{\pi_\theta}(s | s_0) = (1 - \gamma) \sum_{t=0}^{\infty} \gamma^t \mathbb{P}(s_t = s | s_0, \pi_\theta)$ .



香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen

# Policy Gradient Objective

The **policy gradient theorem** states that

$$\begin{aligned}\nabla_{\theta} J(\theta) &\propto \sum_{s \in \mathcal{S}} \rho^{\pi_{\theta}}(s | s_0) \sum_{a \in \mathcal{A}} Q^{\pi_{\theta}}(s, a) \nabla_{\theta} \pi_{\theta}(a | s) \\ &= \mathbb{E}_{\pi} [Q^{\pi_{\theta}}(s, a) \nabla_{\theta} \log \pi_{\theta}(a | s)].\end{aligned}$$



香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen

# Proof of policy gradient theorem

Please refer to the proof of the episodic case in discrete state-action space in the lecture notes.



香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen

# REINFORCE (episodic Monte-Carlo policy-gradient control)

To compute the gradient  $\nabla_{\theta} J(\theta)$  algorithmically, we can sample  $N$  trajectories following the policy  $\pi$  and use the empirical mean to estimate the gradient

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi}[Q^{\pi}(s, a)\nabla_{\theta} \log \pi_{\theta}(a | s)].$$

- For  $Q^{\pi}(s, a)$ , we can use return  $G_t = \sum \gamma^t r_t$  to estimate.
- For  $\nabla_{\theta} \log \pi_{\theta}(a | s)$ , it depends on the form of the policy.



香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen

# REINFORCE (episodic Monte-Carlo policy-gradient control)

---

**Algorithm 1:** REINFORCE (Monte-Carlo method)

---

Initialize the policy parameter  $\theta$   
**for** each episode **do**  
    Sample one trajectory on policy  $\pi_\theta$ :  $s_0, a_0, r_0, s_1, a_1, \dots, s_T$   
    **for** each  $t = 0, 1, \dots, T$  **do**  
         $G_t \leftarrow \sum_{t'=t}^T \gamma^{t'-t} r_{t'}$   
         $\theta \leftarrow \theta + \alpha \gamma^t G_t \nabla_\theta \log \pi_\theta(a_t | s_t)$

---



香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen

## REINFORCE with baselines

One problem of policy gradient method is high variance. (why? [Click to see a very intuitive explanation.](#)) A natural solution is to subtract a baseline  $b(s)$  from  $Q^\pi$ , i.e.,

$$\nabla_\theta J(\theta) \propto \sum_{s \in S} \rho^\pi(s | s_0) \sum_{a \in \mathcal{A}} (Q^\pi(s, a) - b(s)) \nabla \pi_\theta(a | s).$$

The baseline can be any function, even a random variable, as long as it does not depend on the action  $a$ .

$$\sum_a b(s) \nabla \pi(a | s, \theta) = b(s) \nabla \sum_a \pi(a | s, \theta) = b(s) \nabla 1 = 0.$$

The expectation value does not change. The update rule that we end up with is a new version of REINFORCE that includes a general baseline

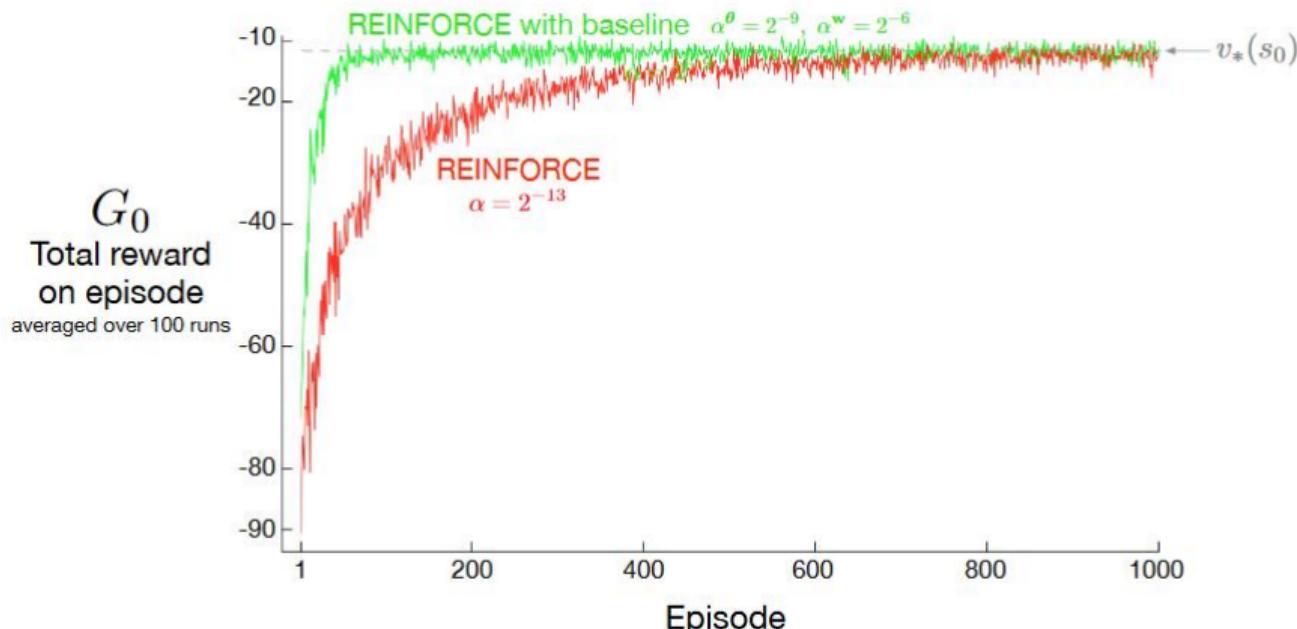


香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen

$$\theta \leftarrow \theta + \alpha \gamma^t (G_t - b(s_t)) \nabla_\theta \log \pi_\theta(a_t | s_t).$$

# REINFORCE with baselines



The Chinese University of Hong Kong, Shenzhen

# REINFORCE with baselines

One natural choice for the baseline is an estimate of the state value  $\hat{V}(s, \mathbf{w})$ , where  $\mathbf{w} \in \mathbb{R}^d$  is a weight vector to be learned. We can use the same method as we adopted in learning  $\theta$  to learn  $\mathbf{w}$ . The complete process is as follows. We have two inputs:

- A differentiable policy parametrization  $\pi_\theta(a | s)$ ;
- A differentiable state value function parametrization  $\hat{V}(s, \mathbf{w})$ .



香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen

# REINFORCE with baselines

---

**Algorithm 2:** REINFORCE with baseline

---

Initialize the policy parameter  $\theta$  and  $w$  at random.

**for** each episode **do**

    Sample one trajectory under policy  $\pi_\theta$ :  $s_0, a_0, r_0, s_1, a_1, r_1 \dots, s_T$

**for** each  $t = 1, 2, \dots, T$  **do**

$$G_t \leftarrow \sum_{t'=t}^T \gamma^{t'-t} r_{t'}$$

$$\delta \leftarrow G_t - \hat{V}(s_t, w)$$

$$w \leftarrow w + \alpha_w \delta \nabla_w \hat{V}(s_t, w)$$

$$\theta \leftarrow \theta + \alpha_\theta \gamma^t \delta \nabla_\theta \log \pi_\theta(a_t | s_t)$$



香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen

# Question and Answering (Q&A)



香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen